

## Logistic Regression and Neural Networks

José Eduardo Ochoa Luna

Dr. Ciencias - Universidade de São Paulo


Maestría C.C. Universidad Católica San Pablo

Sistemas Inteligentes

18 de octubre 2018

# Overfitting

Consider the problem of predicting  $y$  from  $x \in \mathbb{R}$



- 1)  $y = \theta_0 + \theta_1 x$
- 2)  $y = \theta_0 + \theta_1 x + \theta_2 x^2$
- 3)  $y = \sum_{j=0}^5 \theta_j x^j$

# Underfitting and Overfitting

- 1) **underfitting** : data shows structure not captured by the model

# Underfitting and Overfitting

- 1) **underfitting** : data shows structure not captured by the model
- 3) **overfitting**: the fitted curve passes through the data perfectly

# Underfitting and Overfitting

- 1) **underfitting** : data shows structure not captured by the model
- 3) **overfitting**: the fitted curve passes through the data perfectly
- We would not expect this to be a very good predictor of, housing prices ( $y$ ) for different living areas ( $x$ )

# Logistic Regression

Logistic Regression

# Classification

- This is just like the regression problem, except that the values  $y$  take on only a small number of discrete values

# Classification

- This is just like the regression problem, except that the values  $y$  take on only a small number of discrete values
- Focus on the binary classification problem in which  $y$  can take on only two values, 0 and 1.



# Classification

- This is just like the regression problem, except that the values  $y$  take on only a small number of discrete values
- Focus on the binary classification problem in which  $y$  can take on only two values, 0 and 1.
- spam classifier:  $x^{(i)}$  may be some features of a piece of email,  $y$  may be 1 if it is a piece of spam mail, and 0 otherwise

# Classification

- This is just like the regression problem, except that the values  $y$  take on only a small number of discrete values
- Focus on the binary classification problem in which  $y$  can take on only two values, 0 and 1.
- spam classifier:  $x^{(i)}$  may be some features of a piece of email,  $y$  may be 1 if it is a piece of spam mail, and 0 otherwise
- Given  $x^{(i)}$ , the corresponding  $y^{(i)}$  is also called the label for the training example.

# Logistic Regression

- We could use linear regression algorithm

# Logistic Regression

- We could use linear regression algorithm
- It is easy to construct example where this method performs very poorly

# Logistic Regression

- We could use linear regression algorithm
- It is easy to construct example where this method performs very poorly
- It also doesn't make sense for  $h_{\theta}(x)$  to take values larger than 1 or smaller than 0 when we know that  $y \in \{0, 1\}$

# Logistic Regression

- We could use linear regression algorithm
- It is easy to construct example where this method performs very poorly
- It also doesn't make sense for  $h_{\theta}(x)$  to take values larger than 1 or smaller than 0 when we know that  $y \in \{0, 1\}$
- Thus, the hypotheses  $h_{\theta}(x)$  change:

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

# Logistic Regression


- We could use linear regression algorithm
- It is easy to construct example where this method performs very poorly
- It also doesn't make sense for  $h_{\theta}(x)$  to take values larger than 1 or smaller than 0 when we know that  $y \in \{0, 1\}$
- Thus, the hypotheses  $h_{\theta}(x)$  change:

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

- where  $g$  is called the logistic or the sigmoid function

$$g(z) = \frac{1}{1 + e^{-z}}$$


# Sigmoid Function



- $g(z)$  tends towards 1 as  $z \rightarrow \infty$




# Sigmoid Function




- $g(z)$  tends towards 1 as  $z \rightarrow \infty$
- $g(z)$  tends towards 0 as  $z \rightarrow -\infty$

# Sigmoid Function



- $g(z)$  tends towards 1 as  $z \rightarrow \infty$
- $g(z)$  tends towards 0 as  $z \rightarrow -\infty$
- $g(z)$ , and hence  $h(x)$ , is always bounded between 0 and 1

# Sigmoid Function



- $g(z)$  tends towards 1 as  $z \rightarrow \infty$
- $g(z)$  tends towards 0 as  $z \rightarrow -\infty$
- $g(z)$ , and hence  $h(x)$ , is always bounded between 0 and 1
- $x_0 = 1$ , and  $\theta^T x = \theta_0 + \sum_{j=1}^n \theta_j x_j$

# Derivative of the Sigmoid Function

$$\begin{aligned}g'(z) &= \frac{d}{dz} \frac{1}{1 + e^{-z}} \\&= \frac{1}{(1 + e^{-z})^2} (e^{-z}) \\&= \frac{1}{(1 + e^{-z})} \cdot \left(1 - \frac{1}{(1 + e^{-z})}\right) \\&= g(z)(1 - g(z)).\end{aligned}$$

# Likelihood

Let us assume that

$$p(y = 1|x; \theta) = h_{\theta}(x)$$

$$p(y = 0|x; \theta) = 1 - h_{\theta}(x)$$

this can be written more compactly as

$$p(y|x; \theta) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$$

## Likelihood II

Assuming that the  $m$  training examples were generated independently, the likelihood of the parameters is

$$\begin{aligned} L(\theta) &= p(\vec{y}|X; \theta) \\ &= \prod_{i=1}^m p(y^{(i)}|x^{(i)}; \theta) \\ &= \prod_{i=1}^m (h_{\theta}(x^{(i)}))^{(y^{(i)})} (1 - h_{\theta}(x^{(i)}))^{1-y^{(i)}} \end{aligned}$$

It will be easier to maximize the log likelihood:

$$\begin{aligned} l(\theta) &= \log L(\theta) \\ &= \sum_{i=1}^m y^{(i)} \log h(x^{(i)}) + (1 - y^{(i)}) \log(1 - h(x^{(i)})) \end{aligned}$$

# Gradient Ascent

To maximize the likelihood, we can use gradient ascent

$$\theta := \theta + \alpha \nabla_{\theta} l(\theta)$$

Consider on training example  $(x, y)$  and take derivatives to derive stochastic gradient ascent rule:

$$\begin{aligned} \frac{\partial}{\partial(\theta)} l(\theta) &= \left( y \frac{1}{g(\theta^T x)} - (1 - y) \frac{1}{1 - g(\theta^T x)} \right) \frac{\partial}{\partial \theta_j} g(\theta^T x) \\ &= \left( y \frac{1}{g(\theta^T x)} - (1 - y) \frac{1}{1 - g(\theta^T x)} \right) g(\theta^T x)(1 - g(\theta^T x)) \frac{\partial}{\partial \theta_j} \theta^T x \\ &= (y(1 - g(\theta^T x)) - (1 - y)g(\theta^T x)) x_j \\ &= (y - h_{\theta}(x)) x_j \end{aligned}$$

## Gradient Ascent II

We used the fact that  $g'(z) = g(z)(1 - g(z))$ , this gives us the stochastic gradient ascent rule

$$\theta_j := \theta_j + \alpha(y^{(i)} - h_{\theta}(x^{(i)}))x_j^{(i)}$$

- If we compare to the LMS update rule, it looks identical



## Gradient Ascent II

We used the fact that  $g'(z) = g(z)(1 - g(z))$ , this gives us the stochastic gradient ascent rule

$$\theta_j := \theta_j + \alpha(y^{(i)} - h_{\theta}(x^{(i)}))x_j^{(i)}$$

- If we compare to the LMS update rule, it looks identical
- but  $h_{\theta}(x^{(i)})$  is defined as a non-linear function of  $\theta^T x^{(i)}$

# Model Selection

- Given  $h_{\theta}(x) = g(\theta_0 + \theta_1x + \theta_2x^2 + \dots + \theta_kx^k)$

# Model Selection

- Given  $h_{\theta}(x) = g(\theta_0 + \theta_1x + \theta_2x^2 + \dots + \theta_kx^k)$
- We wish to decide if  $k$  should be  $0, 1, \dots$  or  $10$

# Model Selection

- Given  $h_{\theta}(x) = g(\theta_0 + \theta_1x + \theta_2x^2 + \dots + \theta_kx^k)$
- We wish to decide if  $k$  should be 0, 1, ... or 10
- How can we automatically select a model with a good bias and variance tradeoff?

# Model Selection

- Given  $h_{\theta}(x) = g(\theta_0 + \theta_1x + \theta_2x^2 + \dots + \theta_kx^k)$
- We wish to decide if  $k$  should be  $0, 1, \dots$  or  $10$
- How can we automatically select a model with a good bias and variance tradeoff?
- Assume we have a set of models  $M = \{M_1, \dots, M_d\}$  that we are trying to select among.

# First Attempt

Given a training set  $S$

- Train each model  $M_i$  on  $S$ , to get some hypothesis  $h_i$

# First Attempt

Given a training set  $S$

- Train each model  $M_i$  on  $S$ , to get some hypothesis  $h_i$
- Pick the hypotheses with the smallest training error

# First Attempt

Given a training set  $S$

- Train each model  $M_i$  on  $S$ , to get some hypothesis  $h_i$
- Pick the hypotheses with the smallest training error
- This algorithm does not work, why?



# Cross Validation

Given a training set  $S$

- Randomly split  $S$  into  $S_{\text{train}}$  (say, 70% of data) and  $S_{\text{cv}}$  (the remainder 30%).  $S_{\text{cv}}$  is called the hold-out cross validation set.

# Cross Validation

Given a training set  $S$

- Randomly split  $S$  into  $S_{\text{train}}$  (say, 70% of data) and  $S_{\text{cv}}$  (the remainder 30%).  $S_{\text{cv}}$  is called the hold-out cross validation set.
- Train each model  $M_i$  on  $S_{\text{train}}$  only, to get some hypothesis  $h_i$

# Cross Validation

Given a training set  $S$

- Randomly split  $S$  into  $S_{\text{train}}$  (say, 70% of data) and  $S_{\text{cv}}$  (the remainder 30%).  $S_{\text{cv}}$  is called the hold-out cross validation set.
- Train each model  $M_i$  on  $S_{\text{train}}$  only, to get some hypothesis  $h_i$
- Select and output the hypothesis  $h_i$  that had the smallest error  $\hat{E}_{S_{\text{cv}}}(h_i)$  on the hold out cross validation set.

# k-fold Cross Validation

k-fold cross validation holds out less data each time:

- Randomly split  $S$  into  $k$  disjoint subsets of  $m/k$  training examples each  $(S_1, \dots, S_k)$

# k-fold Cross Validation

k-fold cross validation holds out less data each time:

- Randomly split  $S$  into  $k$  disjoint subsets of  $m/k$  training examples each ( $S_1, \dots, S_k$ )
- For each model  $M_i$ , we evaluate it as follows:  
For  $j = 1, \dots, k$ 
  - Train the model  $M_i$  on  $S_1 \cup \dots \cup S_{j-1} \cup S_{j+1} \cup \dots \cup S_k$  to get some hypothesis  $h_{ij}$
  - Test the hypothesis  $h_{ij}$  on  $S_j$ , to get  $\hat{E}_{S_{cv}}(h_{ij})$

The estimated generalization error of the model  $M_i$  is then calculated as the average of the  $\hat{E}_{S_{cv}}(h_{ij})$ 's

# k-fold Cross Validation

k-fold cross validation holds out less data each time:

- Randomly split  $S$  into  $k$  disjoint subsets of  $m/k$  training examples each  $(S_1, \dots, S_k)$
- For each model  $M_i$ , we evaluate it as follows:  
For  $j = 1, \dots, k$ 
  - Train the model  $M_i$  on  $S_1 \cup \dots \cup S_{j-1} \cup S_{j+1} \cup \dots \cup S_k$  to get some hypothesis  $h_{ij}$
  - Test the hypothesis  $h_{ij}$  on  $S_j$ , to get  $\hat{E}_{S_{CV}}(h_{ij})$

The estimated generalization error of the model  $M_i$  is then calculated as the average of the  $\hat{E}_{S_{CV}}(h_{ij})$ 's

- Pick the model  $M_i$  with the lowest estimated generalization error, and retrain that model on the entire training set  $S$ . The resulting hypothesis is the output as our final answer.

## k-fold Cross Validation II

- A typical choice for  $k$  is 10

## k-fold Cross Validation II

- A typical choice for  $k$  is 10
- This procedure may also be more computationally expensive



## k-fold Cross Validation II

- A typical choice for  $k$  is 10
- This procedure may also be more computationally expensive
- In problems with data is really scarce, sometimes  $k = m$

## k-fold Cross Validation II

- A typical choice for  $k$  is 10
- This procedure may also be more computationally expensive
- In problems with data is really scarce, sometimes  $k = m$
- This method is called **leave-one-out cross validation**

# Feature Selection

- Imagine that you have a problem where the number of features  $n$  is very large (perhaps  $n \gg m$ )

# Feature Selection

- Imagine that you have a problem where the number of features  $n$  is very large (perhaps  $n \gg m$ )
- You suspect that there is only a small number of features that are “relevant”

# Feature Selection

- Imagine that you have a problem where the number of features  $n$  is very large (perhaps  $n \gg m$ )
- You suspect that there is only a small number of features that are “relevant”
- You can apply a feature selection algorithm to reduce the number of features

# Feature Selection

- Imagine that you have a problem where the number of features  $n$  is very large (perhaps  $n \gg m$ )
- You suspect that there is only a small number of features that are “relevant”
- You can apply a feature selection algorithm to reduce the number of features
- Given  $n$  features, there are  $2^n$  possible feature subsets

# Feature Selection

- Imagine that you have a problem where the number of features  $n$  is very large (perhaps  $n \gg m$ )
- You suspect that there is only a small number of features that are “relevant”
- You can apply a feature selection algorithm to reduce the number of features
- Given  $n$  features, there are  $2^n$  possible feature subsets
- It is usually too expensive to enumerate over and compare all  $2^n$  models  $\rightarrow$  heuristic search

# Forward Search

- ① Initialize  $\mathcal{F} = \emptyset$
- ② Repeat {
  - (a) For  $i = 1, \dots, n$  if  $i \notin \mathcal{F}$ , let  $\mathcal{F}_i = \mathcal{F} \cup \{i\}$ , and use some version of cross validation to evaluate features  $\mathcal{F}_i$ .
  - (b) Set  $\mathcal{F}$  to be the best feature subset found on step (a).}
- ③ Select and output the best feature subset that was evaluated during the entire search procedure.



## Forward Search II

- The outer loop ends either when  $\mathcal{F} = \{1, \dots, n\}$  , or when  $|\mathcal{F}|$  exceeds a threshold
- The algorithm is one instantiation of **wrapper model feature selection**
- **Backward search** starts off with  $\mathcal{F} = \{1, \dots, n\}$  and repeatedly deletes features one at time until  $\mathcal{F} = \emptyset$
- Wrapper feature selection algorithms often work well, but can be computationally expensive given they need to make many calls to the learning algorithm

# Filter Feature Selection

- The idea is to compute some simple score  $S(i)$  that measures how informative each feature  $x_i$  is about the class labels  $y$
- Then, we simply pick the  $k$  features with the largest scores  $S(i)$
- $S(i)$  could be the correlation between  $x_i$  and  $y$ , as measured on the training data
- It is more common (particularly for discrete-valued features  $x_i$ ) to choose  $S(i)$  to be the **mutual information** between  $x_i$  and  $y$ :

$$M(x_i, y) = \sum_{x_i \in \{0,1\}} \sum_{y \in \{0,1\}} p(x_i, y) \log \frac{p(x_i, y)}{p(x_i)p(y)}$$

## Filter Feature Selection II

- MI can also be expressed as a Kullback-Leibler (KL) divergence  $MI(x_i, y) = KL(p(x_i, y) || p(x_i)p(y))$
- This gives a measure of how different the probability distributions  $p(x_i, y)$  and  $p(x_i)p(y)$  are.
- If  $x_i$  and  $y$  are independent random variables then we would have  $p(x_i, y) = p(x_i)p(y)$ , and the KL-divergence will be zero
- Then  $x_i$  is clearly very non-informative about  $y$  and thus the score  $S(i)$  should be small
- Conversely, if  $x_i$  is very informative about  $y$ , then  $MI(x_i, y)$  would be large

# Neural Networks

# Neural Networks Example

- Recall the housing price prediction problem



# Neural Networks Example

- Recall the housing price prediction problem
- We wish to prevent negative housing prices by setting the absolute minimum price as zero:



# ReLU Function

- Goal:  $f : x \rightarrow y$

# ReLU Function

- Goal:  $f : x \rightarrow y$
- Simplest possible neural networks as single “neuron” in the network where  $f(x) = \max(ax + b, 0)$ , for some coefficients  $a, b$



# ReLU Function

- Goal:  $f : x \rightarrow y$
- Simplest possible neural networks as single “neuron” in the network where  $f(x) = \max(ax + b, 0)$ , for some coefficients  $a, b$
- This function is called a ReLU (rectified linear unit)

# ReLU Function


- Goal:  $f : x \rightarrow y$
- Simplest possible neural networks as single “neuron” in the network where  $f(x) = \max(ax + b, 0)$ , for some coefficients  $a, b$
- This function is called a ReLU (rectified linear unit)
- A more complex neural network may take the single neuron and stack them together such that one neuron passes its output as input into the next neuron

# ReLU Function

- Goal:  $f : x \rightarrow y$
- Simplest possible neural networks as single “neuron” in the network where  $f(x) = \max(ax + b, 0)$ , for some coefficients  $a, b$
- This function is called a ReLU (rectified linear unit)
- A more complex neural network may take the single neuron and stack them together such that one neuron passes its output as input into the next neuron
- That results in a more complex function


# Housing Prediction Revisited

- features: number of bedrooms, zip code, wealth of the neighborhood




# Housing Prediction Revisited

- features: number of bedrooms, zip code, wealth of the neighborhood
- Building neural networks is analogous to Lego bricks: we take individual neurons and stack them together to create complex neural networks



# Housing Prediction Revisited

- features: number of bedrooms, zip code, wealth of the neighborhood
- Building neural networks is analogous to Lego bricks: we take individual neurons and stack them together to create complex neural networks
- Given three derived features (Family size, walkable school quality), the price of the home depends on these features



# Neural Network Learning

- All you need are the input features  $x$  and the output  $y$  while the neural network will figure out everything in the middle by itself

# Neural Network Learning

- All you need are the input features  $x$  and the output  $y$  while the neural network will figure out everything in the middle by itself
- The process of a neural network learning the intermediate features is called **end-to-end learning**



# Neural Network Learning

- All you need are the input features  $x$  and the output  $y$  while the neural network will figure out everything in the middle by itself
- The process of a neural network learning the intermediate features is called **end-to-end learning**
- The input to a NN is a set of input  $x_1, x_2, x_3, x_4$ .

# Neural Network Learning

- All you need are the input features  $x$  and the output  $y$  while the neural network will figure out everything in the middle by itself
- The process of a neural network learning the intermediate features is called **end-to-end learning**
- The input to a NN is a set of input  $x_1, x_2, x_3, x_4$ .
- We connect these four feature to three neurons (called **hidden units**)

# Neural Network Learning

- All you need are the input features  $x$  and the output  $y$  while the neural network will figure out everything in the middle by itself
- The process of a neural network learning the intermediate features is called **end-to-end learning**
- The input to a NN is a set of input  $x_1, x_2, x_3, x_4$ .
- We connect these four feature to three neurons (called **hidden units**)
- Goal NN: automatically determine three relevant features so as to predict the price of a house

# Neural Network Layers

- Black box: it can be difficult to understand the features it has invented

# Neural Network Layers

- Black box: it can be difficult to understand the features it has invented
- input layer (3 units), hidden layer(4 units), output layer (1 unit)

# Neural Network Layers

- Black box: it can be difficult to understand the features it has invented
- input layer (3 units), hidden layer(4 units), output layer (1 unit)
- the first hidden unit requires the input  $x_1, x_2, x_3$  and outputs a value denoted by  $a_1$

# Neural Network Layers

- Black box: it can be difficult to understand the features it has invented
- input layer (3 units), hidden layer(4 units), output layer (1 unit)
- the first hidden unit requires the input  $x_1, x_2, x_3$  and outputs a value denoted by  $a_1$
- $a$  refers to the neuron's **activation** value

# Neural Network Layers

- Black box: it can be difficult to understand the features it has invented
- input layer (3 units), hidden layer(4 units), output layer (1 unit)
- the first hidden unit requires the input  $x_1, x_2, x_3$  and outputs a value denoted by  $a_1$
- $a$  refers to the neuron's **activation** value
- $a_1^{[1]}$  denote the output value of the first hidden unit in the first hidden layer




# Neural Network Layers

- Black box: it can be difficult to understand the features it has invented
- input layer (3 units), hidden layer(4 units), output layer (1 unit)
- the first hidden unit requires the input  $x_1, x_2, x_3$  and outputs a value denoted by  $a_1$
- $a$  refers to the neuron's **activation** value
- $a_1^{[1]}$  denote the output value of the first hidden unit in the first hidden layer
- The input layer is layer 0, the first hidden layer is 1 and the output is layer 2.

# Notation


$$\begin{aligned}x_1 &= a_1^{[0]} \\x_2 &= a_2^{[0]} \\x_3 &= a_3^{[0]}\end{aligned}$$

# Logistic Regression as Neuron




- We can look at at logistic regression  $g(x)$  as a single neuron

# Logistic Regression as Neuron




- We can look at logistic regression  $g(x)$  as a single neuron
- The input is three features:  $x_1, x_2, x_3$ , it outputs an estimated  $y$  value

# Logistic Regression as Neuron



- We can look at logistic regression  $g(x)$  as a single neuron
- The input is three features:  $x_1, x_2, x_3$ , it outputs an estimated  $y$  value
- break  $g(x)$  two computations: 1)  $z = w^T x + b$  and 2)  $a = \sigma(z)$ , where  $\sigma(z) = 1/(1 + e^{-z})$

# Logistic Regression as Neuron



- We can look at logistic regression  $g(x)$  as a single neuron
- The input is three features:  $x_1, x_2, x_3$ , it outputs an estimated  $y$  value
- break  $g(x)$  two computations: 1)  $z = w^T x + b$  and 2)  $a = \sigma(z)$ , where  $\sigma(z) = 1/(1 + e^{-z})$
- notational difference, before  $z = \theta^T x$ , now  $z = w^T x + b$  ( $W$  will denote a matrix)

# Activation functions

$a = g(z)$ , where  $g(z)$  is some activation function:

$$g(z) = \frac{1}{1 + e^{-z}} \quad (\text{sigmoid})$$

$$g(z) = \max(z, 0) \quad (\text{ReLU})$$

$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} \quad (\text{tanh})$$

In general,  $g(z)$  is a non-linear function

# NN computation

- The first hidden unit in the first hidden layer will perform:  
$$z_1^{[1]} = W_1^{[1]T} x + b_1^{[1]} \text{ and } a_1^{[1]} = g(z_1^{[1]})$$



# NN computation

- The first hidden unit in the first hidden layer will perform:  
$$z_1^{[1]} = W_1^{[1]T} x + b_1^{[1]} \text{ and } a_1^{[1]} = g(z_1^{[1]})$$
- $W_1$  is the first row of  $W$ , a matrix of parameters

# NN computation

- The first hidden unit in the first hidden layer will perform:  
$$z_1^{[1]} = W_1^{[1]T} x + b_1^{[1]} \text{ and } a_1^{[1]} = g(z_1^{[1]})$$
- $W_1$  is the first row of  $W$ , a matrix of parameters
- $W_1^{[1]} \in \mathbb{R}^3$  and  $b_1^{[1]} \in \mathbb{R}$

# NN computation

- The first hidden unit in the first hidden layer will perform:  
 $z_1^{[1]} = W_1^{[1]T} x + b_1^{[1]}$  and  $a_1^{[1]} = g(z_1^{[1]})$
- $W_1$  is the first row of  $W$ , a matrix of parameters
- $W_1^{[1]} \in \mathbb{R}^3$  and  $b_1^{[1]} \in \mathbb{R}$
- second and third hidden units (first hidden layer):

$$\begin{aligned} z_2^{[1]} &= W_2^{[1]T} x + b_2^{[1]} \text{ and } a_2^{[1]} = g(z_2^{[1]}) \\ z_3^{[1]} &= W_3^{[1]T} x + b_3^{[1]} \text{ and } a_3^{[1]} = g(z_3^{[1]}) \end{aligned}$$

# NN computation II

The output layer performs:

$$z_1^{[2]} = W_1^{[2]T} a^{[1]} + b_1^{[2]} \text{ and } a_1^{[2]} = g(z_1^{[2]})$$

where  $a^{[1]}$  is the concatenation of all first layer activations:

$$a^{[1]} = \begin{bmatrix} a_1^{[1]} \\ a_2^{[1]} \\ a_3^{[1]} \\ a_4^{[1]} \end{bmatrix}$$

For regression tasks one typically does not apply a non-linear function to  $a_1^{[2]}$

# Vectorization

One must be careful when using for loops. In order to compute hidden unit activations in the first layer, we must compute:

$$\begin{array}{lll} z_1^{[1]} = W_1^{[1]T} x + b_1^{[1]} & \text{and} & a_1^{[1]} = g(z_1^{[1]}) \\ \vdots & & \vdots \\ z_4^{[1]} = W_4^{[1]T} x + b_4^{[1]} & \text{and} & a_4^{[1]} = g(z_4^{[1]}) \end{array}$$

Deep Learning algorithms have high computational requirements. As a result, code will run very slowly if you use for loops

# Vectorizing the Output Computation

$$\underbrace{\begin{bmatrix} z_1^{[1]} \\ \vdots \\ \vdots \\ z_4^{[1]} \end{bmatrix}}_{z^{[1]} \in \mathbb{R}^{4 \times 1}} = \underbrace{\begin{bmatrix} - & W_1^{[1]T} & - \\ - & W_2^{[1]T} & - \\ & \vdots & \\ - & W_4^{[1]T} & - \end{bmatrix}}_{W^{[1]} \in \mathbb{R}^{4 \times 3}} \underbrace{\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}}_{x \in \mathbb{R}^{3 \times 1}} + \underbrace{\begin{bmatrix} b_1^{[1]} \\ b_2^{[1]} \\ \vdots \\ b_4^{[1]} \end{bmatrix}}_{b^{[1]} \in \mathbb{R}^{4 \times 1}}$$

- matrix notation:  $z^{[1]} = W^{[1]}x + b^{[1]}$

# Vectorizing the Output Computation

$$\underbrace{\begin{bmatrix} z_1^{[1]} \\ \vdots \\ z_4^{[1]} \end{bmatrix}}_{z^{[1]} \in \mathbb{R}^{4 \times 1}} = \underbrace{\begin{bmatrix} - & W_1^{[1]T} & - \\ - & W_2^{[1]T} & - \\ & \vdots & \\ - & W_4^{[1]T} & - \end{bmatrix}}_{W^{[1]} \in \mathbb{R}^{4 \times 3}} \underbrace{\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}}_{x \in \mathbb{R}^{3 \times 1}} + \underbrace{\begin{bmatrix} b_1^{[1]} \\ b_2^{[1]} \\ \vdots \\ b_4^{[1]} \end{bmatrix}}_{b^{[1]} \in \mathbb{R}^{4 \times 1}}$$

- matrix notation:  $z^{[1]} = W^{[1]}x + b^{[1]}$
- $a^{[1]} = g(z^{[1]})$ , use vectorized libraries (element-wise operations, e.g. ./ element-wise division Octave)

# Summary Vectorizing

- Given an input  $x \in \mathbb{R}^3$ ,



# Summary Vectorizing

- Given an input  $x \in \mathbb{R}^3$ ,
- We compute the hidden layer's activations:  
 $z^{[1]} = W^{[1]}x + b^{[1]}$  and  $a^{[1]} = g(z^{[1]})$

# Summary Vectorizing

- Given an input  $x \in \mathbb{R}^3$ ,
- We compute the hidden layer's activations:  
 $z^{[1]} = W^{[1]}x + b^{[1]}$  and  $a^{[1]} = g(z^{[1]})$
- To compute the output layer's activations:

$$\underbrace{z^{[2]}}_{1 \times 1} = \underbrace{W^{[2]}}_{1 \times 4} \underbrace{a^{[1]}}_{4 \times 1} + \underbrace{b^{[2]}}_{1 \times 1} \quad \text{and} \quad \underbrace{a^{[2]}}_{1 \times 1} = g(\underbrace{z^{[2]}}_{1 \times 1})$$

# Non-linear Activation Functions

Why not use  $g(z) = z$ ?. Assume that  $b^{[1]}$  and  $b^{[2]}$  are zeros:

$$\begin{aligned} z^{[2]} &= W^{[2]} a^{[1]} \\ &= W^{[2]} g(z^{[1]}) \\ &= W^{[2]} z^{[1]} \\ &= W^{[2]} W^{[1]} x \\ &= \hat{W} x \end{aligned}$$

- Applying a linear function to another linear function will result in a linear function over the original input
- Without non-linear activation functions, the NN will simply perform linear regression

# Vectorization Over Training Examples

Training set with three examples:

$$\begin{aligned}z^{[1](1)} &= W^{[1]}x^{(1)} + b^{[1]} \\z^{[1](2)} &= W^{[1]}x^{(2)} + b^{[1]} \\z^{[1](3)} &= W^{[1]}x^{(3)} + b^{[1]}\end{aligned}$$

Vectorization

$$X = \begin{bmatrix} \begin{array}{c} | \\ x^{(1)} \\ | \end{array} & \begin{array}{c} | \\ x^{(2)} \\ | \end{array} & \begin{array}{c} | \\ x^{(3)} \\ | \end{array} \end{bmatrix}$$

$$Z^{[1]} = \begin{bmatrix} \begin{array}{c} | \\ z^{[1](1)} \\ | \end{array} & \begin{array}{c} | \\ z^{[1](2)} \\ | \end{array} & \begin{array}{c} | \\ z^{[1](3)} \\ | \end{array} \end{bmatrix} = W^{[1]}X + b^{[1]}$$

# Putting it together

Suppose we have a training set  $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$ , where  $x^{(i)}$  is a picture and  $y^{(i)}$  is a binary label for whether the picture contains a cat or not

- Initialize the parameters  $W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}$  to small random numbers

# Putting it together

Suppose we have a training set  $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$ , where  $x^{(i)}$  is a picture and  $y^{(i)}$  is a binary label for whether the picture contains a cat or not

- Initialize the parameters  $W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}$  to small random numbers
- For each example, we compute the output “probability” from the sigmoid function  $a^{[2](i)}$

# Putting it together

Suppose we have a training set  $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$ , where  $x^{(i)}$  is a picture and  $y^{(i)}$  is a binary label for whether the picture contains a cat or not

- Initialize the parameters  $W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}$  to small random numbers
- For each example, we compute the output “probability” from the sigmoid function  $a^{[2](i)}$
- Using the logistic regression log likelihood:  
 $\sum_{i=1}^m (y^{(i)} \log a^{[2](i)} + (1 - y^{(i)}) \log(1 - a^{[2](i)}))$ , we maximize this function using gradient ascent

# Putting it together

Suppose we have a training set  $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$ , where  $x^{(i)}$  is a picture and  $y^{(i)}$  is a binary label for whether the picture contains a cat or not

- Initialize the parameters  $W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}$  to small random numbers
- For each example, we compute the output “probability” from the sigmoid function  $a^{[2](i)}$
- Using the logistic regression log likelihood:  
$$\sum_{i=1}^m (y^{(i)} \log a^{[2](i)} + (1 - y^{(i)}) \log(1 - a^{[2](i)})),$$
 we maximize this function using gradient ascent
- This maximization procedure corresponds to training the neural network



# Tarea

Implementar ejercicios indicados en:  
logisticRegressionSemana2.ipynb